


Evaluation logic from leakage control to governed clinical use

(a) Leakage-resistant design



(b) Segmentation operators




Report a task-specific bundle—not a single score

(c) Diagnostic / prognostic claims

score + observed outcome + intended case spectrum




prognostic horizon and censoring




Discrimination $\neq$ calibration $\neq$ utility

(d) Aggregation / uncertainty



Confidence $\neq$ correctness under shift

(e) Clinical evidence / governance



Metric adequacy $\neq$ evidence strength $\neq$ clinical utility